

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT
on

BIG DATA ANALYTICS **(20CS6PEBDA)**

Submitted by

KRISHNA MOHAN DULLOLLI (1BM19CS075)

in partial fulfillment for the award of the degree of
BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING
(Autonomous Institution under VTU)
BENGALURU-560019
May-2022 to July-2022

B. M. S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “**BIG DATA ANALYTICS**” carried out by **KRISHNA MOHAN DULLOLLI (1BM19CS075)**, who is bonafide student of **B. M. S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the year 2022. The Lab report has been approved as it satisfies the academic requirements in respect of a **Course Title - (20CS6PEBDA)** work prescribed for the said degree.

Name of the Lab-Incharge
Designation
Department of CSE
BMSCE, Bengaluru

Prof. Pallavi G B
Assistant Professor
Department of CSE
BMSCE, Bengaluru

Index Sheet

Sl. No.	Experiment Title	Page No.
1	MongoDB CRUD Operations	4
2	MongoDB Operations	7
3	Cassandra Lab 1	10
4	Cassandra Lab 2	11
5	Screenshot of Hadoop	12
6	Execution of HDFS	13
7	MapReduce to find mean/max temperature	15
8	List top 10 max occurrence	19
9	Demonstrate Join Operation	21
10	Print word count on scala	24
11	Using RDD and FlatMap	25

Course Outcome

CO1	Apply the concept of NoSQL, Hadoop or Spark for a given task
CO2	Analyze the Big Data and obtain insight using data analytics mechanisms.
CO3	Design and implement Big data applications by applying NoSQL, Hadoop or Spark

1 MongoDB CRUD Operations

I. CREATE DATABASE IN MONGODB

>use KrishnaDB

switched to db KrishnaDB

II. CRUD (CREATE, READ, UPDATE, DELETE) OPERATIONS

>db.createCollection("Student");

{ "ok" : 1 }

>db.Student.insert({_id:1,name:"Krishna",grade:9});

WriteResult({ "nInserted" : 1 })

>db.Student.update({_id:6,name:"qwert"},{\$set:{grade:4}},{upsert:true});

WriteResult({ "nMatched" : 0, "nUpserted" : 1, "nModified" : 0, "_id" : 6 })

>db.Student.find();

{ "_id" : 1, "name" : "Krishna", "grade" : 9 }

{ "_id" : 2, "name" : "Abc", "grade" : 10 }

{ "_id" : 3, "name" : "Mno", "grade" : 5 }

{ "_id" : 4, "name" : "Pqr", "grade" : 8 }

> show collections;

Student

III. Save Method

> db.Student.save({name:"zzz",_id:10,grade:8});

WriteResult({ "nMatched" : 0, "nUpserted" : 1, "nModified" : 0, "_id" : 10 })

IV. COUNT

> db.Student.count();

6

> db.Student.count({grade:9});

1

V FIND

> db.Student.find({grade:{\$lt:5}},{name:1,grade:1,_id:0});

```
{ "grade" : 2, "name" : "qwert" }
```

> db.Student.find({name:{\$in:["Krishna","Abc","Mno"]}}, {name:1,grade:1,_id:0});

```
{ "name" : "Krishna", "grade" : 9 }
```

```
{ "name" : "Abc", "grade" : 10 }
```

```
{ "name" : "Mno", "grade" : 5 }
```

> db.Student.find({name:/^S/},{name:1,grade:1,_id:0});

```
{ "name" : "Krishna", "grade" : 9 }
```

> db.Student.find({name:/.b/},{name:1,grade:1,_id:0});

```
{ "name" : "Abc", "grade" : 10 }
```

> db.Student.find().sort({name:1});

```
{ "_id" : 2, "name" : "Abc", "grade" : 10 }
```

```
{ "_id" : 3, "name" : "Mno", "grade" : 5 }
```

```
{ "_id" : 4, "name" : "Pqr", "grade" : 8 }
```

```
{ "_id" : 1, "name" : "Krishna", "grade" : 9 }
```

```
{ "_id" : 7, "name" : "kkk", "grade" : 6 }
```

```
{ "_id" : 6, "grade" : 2, "name" : "qwert" }
```

> db.Student.find().sort({name:1,grade:-1});

```
{ "_id" : 2, "name" : "Abc", "grade" : 10 }
```

```
{ "_id" : 3, "name" : "Mno", "grade" : 5 }
```

```
{ "_id" : 4, "name" : "Pqr", "grade" : 8 }
```

```
{ "_id" : 1, "name" : "Krishna", "grade" : 9 }
```

```
{ "_id" : 7, "name" : "kkk", "grade" : 6 }
```

```
{ "_id" : 6, "grade" : 2, "name" : "qwert" }
```

> db.Student.find({grade:8}).limit(3);

```
{ "_id" : 4, "name" : "Pqr", "grade" : 8 }
```

```
{ "_id" : 10, "name" : "zzz", "grade" : 8 }
```

```
> db.Student.find().skip(2);
```

```
{ "_id" : 3, "name" : "Mno", "grade" : 5 }
```

```
{ "_id" : 4, "name" : "Pqr", "grade" : 8 }
```

```
{ "_id" : 6, "grade" : 2, "name" : "qwert" }
```

```
{ "_id" : 7, "name" : "kkk", "grade" : 6 }
```

```
{ "_id" : 10, "name" : "zzz", "grade" : 8 }
```

VI. AGGREGATE FUNCTIONS

```
> db.faculty.aggregate ( { $match:{department:"mech"}}, { $group : { _id :  
"$designation", AverageSal : { $avg : "$salary" } } },  
{ $match:{AverageSal:{ $gt:50000 }}});
```

```
{ "_id" : " associate prof", "AverageSal" : 85000 }
```

```
{ "_id" : "assistant prof", "AverageSal" : 70000 }
```

VII. ARRAYS

```
> db.food.insert({_id:1,fruits:['apple','mango']});
```

```
WriteResult({ "nInserted" : 1 })
```

```
> db.food.find({fruits:['pineapple','mango','orange']});
```

```
{ "_id" : 3, "fruits" : [ "pineapple", "mango", "orange" ] }
```

```
> db.food.find({fruits:{ $all:['pineapple']}});
```

```
{ "_id" : 2, "fruits" : [ "pineapple", "mango", "grapes" ] }
```

```
{ "_id" : 3, "fruits" : [ "pineapple", "mango", "orange" ] }
```

```
> db.food.update({_id:2},{ $set:{'fruits.1':'apple'}});
```

```
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

```
> db.food.update({_id:2},{ $push:{price:{grapes:80,mango:200,cherry:100}}} );
```

```
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

2. MongoDB Operations

1) Faculty DB

i) Create a database for Faculty and Create a Faculty Collection(Faculty_id, Name, Designation ,Department, Age, Salary, Specialization(Set)).

>use Faculty

> db.createCollection("faculty")

ii) Insert required documents to the collection.

> db.faculty.insert({_id:1,name:"abc",designation:"assistant prof",department:"mech",age:31,salary:90000,specialization:['python','mysql','autocad']});

iii) First Filter on “Dept_Name:MECH” and then group it on “Designation” and compute the Average Salary for that Designation and filter those documents where the “Avg_Sal” is greater than 650000.

> db.faculty.aggregate ({\$match:{department:"mech"}}, {\$group : {_id : "\$designation", AverageSal :{\$avg:"\$salary"} } }, {\$match:{AverageSal:{\$gt:50000}}});

```
{ "_id" : " associate prof", "AverageSal" : 85000 }
```

```
{ "_id" : "assistant prof", "AverageSal" : 70000 }
```

2) Consider a table “Product” with the following columns:

Product _id

ProductName

ManufacturingDate

Price

Quantity

Write MongoDB queries for the following:

> use Products switched to db Products

> db.createCollection("product");

```
{ "ok" : 1 }
```

```
>
db.product.insert({pid:1,pname:"keyboard",mdate:2001,price:1800,quantity:2})
;
```

```
WriteResult({ "nInserted" : 1 })
```

i)To display only the product name from all the documents of the product collection.

```
> db.product.find({}, {pname:1, _id:0});
```

```
{ "pname" : "keyboard" }
```

```
{ "pname" : "mouse" }
```

```
{ "pname" : "motherboard" }
```

ii)To display only the Product ID, ExpiryDate as well as the quantity from the document of the product collection where the _id column is 1.

```
> db.product.find({pid:1},
```

```
{pid:1, _id:0, mdate:1, quantity:1});
```

```
{ "pid" : 1, "mdate" : 2001, "quantity" : 2 }
```

iii) To find those documents where the price is not set to 45000.

```
> db.product.find({price:{$ne:45000}}, {pname:1, _id:0});
```

```
{ "pname" : "keyboard" }
```

```
{ "pname" : "mouse" }
```

```
{ "pname" : "motherboard" }
```

iv)To find those documents from the Product collection where the quantity is set to 30 and the product name is set to 'LEDTV'.

```
> db.product.find({$and:[{quantity:{$eq:30}}, {pname:{$eq:"LED
TV"}}]}, {pname:1, _id:0})8
```

```
{ "pname" : "LED TV" }
```

v)To find documents from the Product collection where the Product name ends in 'r'.

```
> db.product.find({pname:/d$/}, {pname:1, quantity:1, _id:0})
```

```
{ "pname" : "keyboard", "quantity" : 2 }
```



```
{ "pname" : "motherboard", "quantity" : 150 }
```

3) Create a mongodb collection Hospital. Demonstrate the following by choosing fields of your choice.

> use Hospital switched to db Hospital

> db.createCollection("hospital");

```
{ "ok" : 1 }
```

> db.hospital.insert({_id:1,name:"xyz",diseases:["diabetes","high bp","fever"]});

```
WriteResult({ "nInserted" : 1 })
```

1. Insert three documents

> db.hospital.updateMany({},{\$pull:{diseases:"fever"}});

```
{ "acknowledged" : true, "matchedCount" : 3, "modifiedCount" : 2 }
```

2. Use Arrays(Use Pull and Pop operation)

> db.hospital.updateOne({_id:1},{\$pop:{diseases:-1}});

```
{ "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
```

3. Use Index

> db.hospital.find({"diseases.2":"nausea"});

```
{ "_id" : 3, "name" : "mno", "diseases" : [ "covid", "sarscov", "nausea" ] }
```

4. Use Cursors

> db.hospital.find({}).count();

```
3
```

> db.hospital.find({}).limit(2);

```
{ "_id" : 1, "name" : "xyz", "diseases" : [ "high bp" ] } { "_id" : 2, "name" : "abc", "diseases" : [ "typhoid", "cholera" ] }
```

> db.hospital.find({}).size();

```
3
```

5. Updation

> db.hospital.update({_id:3},{\$set:{'diseases.1':'sarscov'}});

```
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

3. Cassandra Lab 1

1. Create a key space by name Employee

```
cqlsh:saf> create keyspace Employee with  
replication={'class':'SimpleStrategy','replication_factor':1}; cqlsh:saf> use  
Employee ;
```

2. Create a column family by name Employee-Info with attributes Emp_Id Primary Key, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name

```
cqlsh:employee> create table empInfo( emp_id int PRIMARY KEY, emp_name  
text,desig text,dpj timestamp,salary int,dept_name text );
```

3. Insert the values into the table in batch

```
cqlsh:employee> insert into  
empInfo(emp_id,emp_name,desig,dpj,salary,dept_name) values( 1,  
'Krishna','sde', '2022-05-05', 200000, 'cse' );
```

4. Update Employee name and Department of Emp-Id 121

```
cqlsh:employee> update empInfo set emp_name='zzz',dept_name='ie'where  
emp_id=2;
```

5. Sort the details of Employee records based on salary

```
.cqlsh:employee> select * from emp_Info where emp_id in (1,2,3) order by  
salary;
```

6. Alter the schema of the table Employee_Info to add a column Projects;which stores a set of Projects done by the corresponding Employee.

```
cqlsh:employee> alter table empInfo add project set
```

7. Update the altered table to add project names.

```
cqlsh:employee> update empInfo set project={'reactJs','MI'} where emp_id=1;
```

8 Create a TTL of 15 seconds to display the values of Employees.

```
cqlsh:employee> insert into  
empInfo(emp_id,emp_name,desig,dpj,salary,dept_name) values( 5, 'wxy', 'sde',  
'2022-02-05', 250000, 'cse' ) using ttl 30; cqlsh:employee> select ttl(emp_name)  
from empInfo;
```

4. Cassandra Lab 2

1 Create a key space by name Library

```
CREATE keyspace library1 with replication={ 'class':'SimpleStrategy',  
'replication_factor':1 };
```

2. Create a column family by name Library-Info with attributes Stud_Id Primary Key,Counter_value of type Counter,Stud_Name, Book-Name, Book-Id, Date_of_issue

```
CREATE TABLE lib.libinfo1 ( s_id int, sname text, book text, bid int, doi  
timestamp, counter_val counter, PRIMARY KEY (s_id, sname, book, bid, doi) );
```

3. Insert the values into the table in batch

```
update libinfo set counter_val=counter_val+1 where s_id=1 and sname='saf'  
and book='harry potter1' and bid=1 and doi='2022-05-05';
```

4. Display the details of the table created and increase the value of the counter

```
cqlsh:lib> update libinfo set counter_val=counter_val+1 where s_id=1 and  
sname='saf' and book='harry potter1'; cqlsh:lib> select * from libinfo;
```

5. Write a query to show that a student with id 112 has taken a book “BDA” 2 times.

```
cqlsh:lib> select counter_val from libinfo where s_id=1 and sname='saf' and  
book='harry potter1';
```

```
counter_val
```

```
2
```

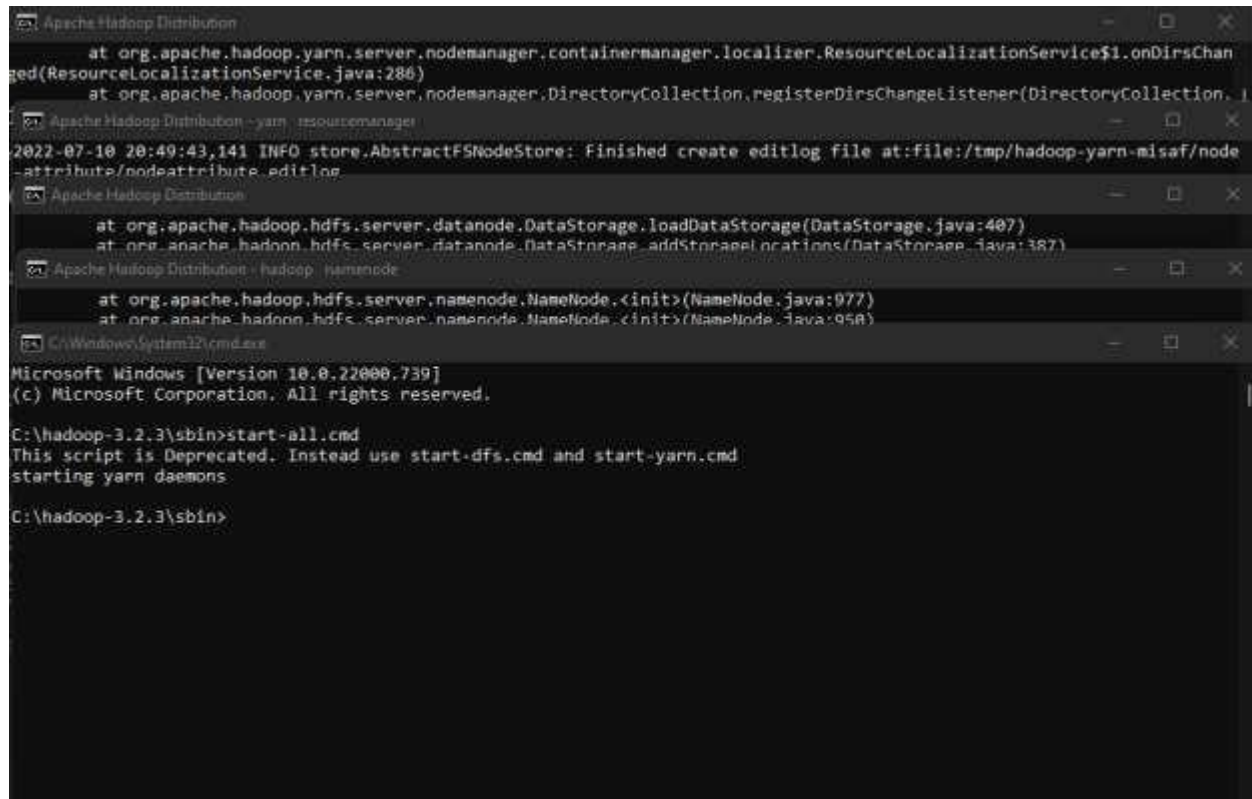
6. Export the created column to a csv file

```
COPY libinfo(s_id,sname,book,bid,doi,counter_val) TO 'data1.csv' WITH HEADER  
= TRUE;
```

7. Import a given csv dataset from local file system into Cassandra column family

```
COPY libinfo(s_id,sname,book,bid,doi) FROM 'libdata.csv' WITH HEADER =  
TRUE;
```

5. Screenshot of Hadoop Installation



The screenshot displays a series of overlapping windows from a Hadoop installation. The top window, titled 'Apache Hadoop Distribution', shows Java stack traces for 'ResourceLocalizationService' and 'DirectoryCollection'. Below it, a window titled 'Apache Hadoop Distribution - yarn - resourcemanager' shows a log message: '2022-07-10 20:49:43,141 INFO store.AbstractFSNodeStore: Finished create editlog file at:file:/tmp/hadoop-yarn-misaf/node-attribute/nodeattribute.editlog'. Another window titled 'Apache Hadoop Distribution' shows stack traces for 'DataStorage' and 'NameNode'. The bottom window is a Windows command prompt titled 'C:\Windows\System32\cmd.exe', showing the execution of 'start-all.cmd' with a deprecation warning and the current directory 'C:\hadoop-3.2.3\sbin>'.

```
at org.apache.hadoop.yarn.server.nodemanager.containermanager.localizer.ResourceLocalizationService$1.onDirChange
ged(ResourceLocalizationService.java:286)
at org.apache.hadoop.yarn.server.nodemanager.DirectoryCollection.registerDirChangeListener(DirectoryCollection, 1
Apache Hadoop Distribution - yarn - resourcemanager
2022-07-10 20:49:43,141 INFO store.AbstractFSNodeStore: Finished create editlog file at:file:/tmp/hadoop-yarn-misaf/node
-attribute/nodeattribute.editlog
Apache Hadoop Distribution
at org.apache.hadoop.hdfs.server.datanode.DataStorage.loadDataStorage(DataStorage.java:407)
at org.apache.hadoop.hdfs.server.datanode.DataStorage.addStorageLocations(DataStorage.java:387)
Apache Hadoop Distribution - hadoop - namenode
at org.apache.hadoop.hdfs.server.namenode.NameNode.<init>(NameNode.java:977)
at org.apache.hadoop.hdfs.server.namenode.NameNode.<init>(NameNode.java:958)
C:\Windows\System32\cmd.exe
Microsoft Windows [Version 10.0.22000.739]
(c) Microsoft Corporation. All rights reserved.

C:\hadoop-3.2.3\sbin>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons

C:\hadoop-3.2.3\sbin>
```

6. HDFS Commands

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /
21/04/19 23:41:08 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:19 /mydir
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:21 /mydr
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:39 /newdir
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -mv /mydr /newdir
21/04/19 23:41:38 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /
21/04/19 23:41:44 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:19 /mydir
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:41 /newdir
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /newdir
21/04/19 23:42:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:21 /newdir/mydr
hduser@lab-VirtualBox:/usr/local/sbin$
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /
21/04/19 23:52:26 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:45 /mydir
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:48 /newdir
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -rm -R /mydir
21/04/19 23:52:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
21/04/19 23:52:57 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
Deleted /mydir
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /
21/04/19 23:53:02 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
drwxr-xr-x  - hduser supergroup          0 2021-04-19 23:48 /newdir
hduser@lab-VirtualBox:/usr/local/sbin$
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -mkdir /mydir
21/04/19 22:58:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /
21/04/19 22:58:36 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x  - hduser supergroup          0 2021-04-19 22:58 /mydir
drwxr-xr-x  - hduser supergroup          0 2021-04-18 19:27 /mydr
```



```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -get /mydir ~/copyfromhadoop
21/04/19 23:25:49 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /
21/04/19 23:48:41 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x   - hduser supergroup          0 2021-04-19 23:45 /mydir
drwxr-xr-x   - hduser supergroup          0 2021-04-19 23:41 /newdir
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -cp /mydir/sample.txt /newdir
21/04/19 23:48:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /newdir
21/04/19 23:49:22 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x   - hduser supergroup          0 2021-04-19 23:21 /newdir/mydir
-rw-r--r--   1 hduser supergroup         13 2021-04-19 23:48 /newdir/sample.txt
hduser@lab-VirtualBox:/usr/local/sbin$
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -copyToLocal /mydir ~/hadoopcopy
21/04/19 23:29:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hduser@lab-VirtualBox:/usr/local/sbin$
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -copyFromLocal ~/file1.txt /mydir
21/04/19 23:19:36 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -ls /mydir
21/04/19 23:20:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r--   1 hduser supergroup         30 2021-04-19 23:19 /mydir/file1.txt
hduser@lab-VirtualBox:/usr/local/sbin$
```

```
hduser@lab-VirtualBox:/usr/local/sbin$ hadoop fs -cat /mydir/file1.txt
21/04/19 23:38:07 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
I am using Hadoop
line1
line2
hduser@lab-VirtualBox:/usr/local/sbin$
```

7. Mean/Max temperature of weather data

Driver class:

```
package temperatureMax;

import org.apache.hadoop.io.*;
import org.apache.hadoop.fs.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class TempDriver
{
    public static void main (String[] args) throws Exception
    {
        if (args.length != 2)
        {
            System.err.println("Please Enter the input and output parameters");
            System.exit(-1);
        }
        Job job = new Job();
        job.setJarByClass(TempDriver.class);
        job.setJobName("Max temperature");
        FileInputFormat.addInputPath(job,new Path(args[0]));
        FileOutputFormat.setOutputPath(job,new Path (args[1]));
        job.setMapperClass(TempMapper.class);
        job.setReducerClass(TempReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        System.exit(job.waitForCompletion(true)?0:1);
    }
}
```

Mapper Class

```
package temperatureMax;

import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import java.io.IOException;

public class TempMapper extends Mapper <LongWritable, Text, Text, IntWritable>
{
    public static final int MISSING = 9999;

    public void map(LongWritable key, Text value, Context context) throws IOException,
    InterruptedException
    {
        String line = value.toString();
        String month = line.substring(19,21);
        int temperature;
        if (line.charAt(87)=='+')
            temperature = Integer.parseInt(line.substring(88, 92));
        else
            temperature = Integer.parseInt(line.substring(87, 92));
        String quality = line.substring(92, 93);
        if(temperature != MISSING && quality.matches("[01459]"))
            context.write(new Text(month),new IntWritable(temperature));
    }
}
```


Reducer class

```
package temperatureMax;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.*;
import java.io.IOException;

public class TempReducer extends Reducer <Text, IntWritable,Text, IntWritable>
{
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws
    IOException,InterruptedException
    {
        int max_temp = 0;
        for (IntWritable value : values)
        {
            if(max_temp<value.get()) {
                max_temp = value.get();
            }
        }
        context.write(key, new IntWritable(max_temp));
    }
}
```

Output:

```
hduser@lab-VirtualBox:/home/lab$ hadoop dfs -cat /tempmax/part-r-00000
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication
.util.KerberosUtil (file:/usr/local/hadoop/share/hadoop/common/lib/hadoop-auth-
2.6.0.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop
.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflec
tive access operations
WARNING: All illegal access operations will be denied in a future release
21/05/10 16:08:48 WARN util.NativeCodeLoader: Unable to load native-hadoop libr
ary for your platform... using builtin-java classes where applicable
03      111
05      22
```

```
hduser@lab-VirtualBox:/home/lab$ hadoop dfs -ls /tempmax
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication
.util.KerberosUtil (file:/usr/local/hadoop/share/hadoop/common/lib/hadoop-auth-
2.6.0.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop
.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflec
tive access operations
WARNING: All illegal access operations will be denied in a future release
21/05/10 16:08:23 WARN util.NativeCodeLoader: Unable to load native-hadoop libr
ary for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--    1 hduser supergroup          0 2021-05-10 16:08 /tempmax/_SUCCESS
-rw-r--r--    1 hduser supergroup       13 2021-05-10 16:08 /tempmax/part-r-00
000
```

8. Word Occurences

```
hduser@lab-VirtualBox:/home/lab/hadoop-2.6.0/share/hadoop/mapreduce$ hadoop jar
hadoop-mapreduce-examples-2.6.0.jar wordcount /input /firstExampleOut
21/04/26 15:19:29 WARN util.NativeCodeLoader: Unable to load native-hadoop libr
ary for your platform... using builtin-java classes where applicable
21/04/26 15:19:31 INFO Configuration.deprecation: session.id is deprecated. Ins
tead, use dfs.metrics.session-id
21/04/26 15:19:31 INFO jvm.JvmMetrics: Initializing JVM Metrics with processNam
e=JobTracker, sessionId=
21/04/26 15:19:32 INFO input.FileInputFormat: Total input paths to process : 1
21/04/26 15:19:32 INFO mapreduce.JobSubmitter: number of splits:1
21/04/26 15:19:33 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_l
ocal1167934544_0001
21/04/26 15:19:33 INFO mapreduce.Job: The url to track the job: http://localhos
t:8080/
21/04/26 15:19:33 INFO mapreduce.Job: Running job: job_local1167934544_0001
21/04/26 15:19:33 INFO mapred.LocalJobRunner: OutputCommitter set in config nul
l
21/04/26 15:19:33 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.had
oop.mapreduce.lib.output.FileOutputCommitter
21/04/26 15:19:34 INFO mapreduce.Job: Job job_local1167934544_0001 running in u
ber mode : false
21/04/26 15:19:34 INFO mapreduce.Job: map 0% reduce 0%
21/04/26 15:19:34 INFO mapred.LocalJobRunner: Waiting for map tasks
21/04/26 15:19:34 INFO mapred.LocalJobRunner: Starting task: attempt_local11679
34544_0001_m_000000_0
21/04/26 15:19:34 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
21/04/26 15:19:34 INFO mapred.MapTask: Processing split: hdfs://localhost:54310
/input/input.txt:0+44
```

```
21/04/26 15:19:37 INFO mapreduce.Job: map 100% reduce 100%
21/04/26 15:19:37 INFO mapreduce.Job: Job job_local1167934544_0001 completed
successfully
21/04/26 15:19:37 INFO mapreduce.Job: Counters: 38
File System Counters
  FILE: Number of bytes read=541122
  FILE: Number of bytes written=1047373
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=88
  HDFS: Number of bytes written=35
  HDFS: Number of read operations=15
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=4
Map-Reduce Framework
  Map input records=3
  Map output records=9
  Map output bytes=80
  Map output materialized bytes=61
  Input split bytes=103
  Combine input records=9
  Combine output records=5
  Reduce input groups=5
  Reduce shuffle bytes=61
  Reduce input records=5
  Reduce output records=5
  Spilled Records=10
  Shuffled Maps =1
  Failed Shuffles=0
```



```
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=340787200

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
Bytes Read=44
File Output Format Counters
Bytes Written=35

hduser@lab-VirtualBox: /home/lab/hadoop-2.6.0/share/hadoop/mapreduce$ hadoop dfs
-ls /firstExampleOut
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

21/04/26 15:20:15 WARN util.NativeCodeLoader: Unable to load native-hadoop libr
ary for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  1 hduser supergroup          0 2021-04-26 15:19 /firstExampleOut/_
SUCCESS
-rw-r--r--  1 hduser supergroup        35 2021-04-26 15:19 /firstExampleOut/p
art-r-000000
hduser@lab-VirtualBox: /home/lab/hadoop-2.6.0/share/hadoop/mapreduce$ hadoop dfs
-cat /firstExampleOut/part-r-000000
DEPRECATED: Use of this script to execute hdfs command is deprecated.
```

```
Bytes Read=44
File Output Format Counters
Bytes Written=35

hduser@lab-VirtualBox: /home/lab/hadoop-2.6.0/share/hadoop/mapreduce$ hadoop dfs
-ls /firstExampleOut
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

21/04/26 15:20:15 WARN util.NativeCodeLoader: Unable to load native-hadoop libr
ary for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  1 hduser supergroup          0 2021-04-26 15:19 /firstExampleOut/_
SUCCESS
-rw-r--r--  1 hduser supergroup        35 2021-04-26 15:19 /firstExampleOut/p
art-r-000000
hduser@lab-VirtualBox: /home/lab/hadoop-2.6.0/share/hadoop/mapreduce$ hadoop dfs
-cat /firstExampleOut/part-r-000000
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

21/04/26 15:22:01 WARN util.NativeCodeLoader: Unable to load native-hadoop libr
ary for your platform... using builtin-java classes where applicable
bear      2
car       3
deer      1
deer      1
river     2
hduser@lab-VirtualBox: /home/lab/hadoop-2.6.0/share/hadoop/mapreduce$ hadoop dfs
-cat /firstExampleOut/part-r-000000
```

9. Use of Join

Driver Class

```
package MapReduceJoin;

import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.mapred.lib.MultipleInputs;
import org.apache.hadoop.util.*;

public class JoinDriver extends Configured implements Tool {

    public static class KeyPartitioner implements Partitioner<TextPair, Text> {
        @Override
        public void configure(JobConf job) {}

        @Override
        public int getPartition(TextPair key, Text value, int numPartitions) {
            return (key.getFirst().hashCode() & Integer.MAX_VALUE) % numPartitions;
        }
    }

    @Override
    public int run(String[] args) throws Exception {

        if (args.length != 3) {
            System.out.println("Usage: <Department Emp Strength input> <Department  
Name input> <output>");
            return -1;
        }

        JobConf conf = new JobConf(getConf(), getClass());
        conf.setJobName("Join 'Department Emp Strength input' with 'Department Name  
input'");

        Path AInputPath = new Path(args[0]);
        Path BInputPath = new Path(args[1]);
        Path outputPath = new Path(args[2]);

        MultipleInputs.addInputPath(conf, AInputPath, TextInputFormat.class,
DeptNameMapper.class);
        MultipleInputs.addInputPath(conf, BInputPath, TextInputFormat.class,
DeptEmpStrengthMapper.class);
    }
}
```

```

        FileOutputFormat.setOutputPath(conf, outputPath);

        conf.setPartitionerClass(KeyPartitioner.class);
        conf.setOutputValueGroupingComparator(TextPair.FirstComparator.class);

        conf.setMapOutputKeyClass(TextPair.class);

        conf.setReducerClass(JoinReducer.class);

        conf.setOutputKeyClass(Text.class);

        JobClient.runJob(conf);

        return 0;
    }

    public static void main(String[] args) throws Exception {

        int exitCode = ToolRunner.run(new JoinDriver(), args);
        System.exit(exitCode);
    }
}

```

Mapper Class

```

package MapReduceJoin;

import java.io.IOException;

import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class DeptNameMapper extends MapReduceBase implements Mapper<LongWritable, Text,
TextPair, Text> {

    @Override
    public void map(LongWritable key, Text value, OutputCollector<TextPair, Text> output, Reporter
reporter)
        throws IOException
    {
        String valueString = value.toString();
        String[] SingleNodeData = valueString.split("\t");
        output.collect(new TextPair(SingleNodeData[0], "0"), new Text(SingleNodeData[1]));
    }
}

```

Reducer Class

```
package MapReduceJoin;

import java.io.IOException;
import java.util.Iterator;

import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class JoinReducer extends MapReduceBase implements Reducer<TextPair, Text, Text, Text> {

    @Override
    public void reduce (TextPair key, Iterator<Text> values, OutputCollector<Text, Text> output,
Reporter reporter)
        throws IOException
    {

        Text nodeId = new Text(values.next());
        while (values.hasNext()) {
            Text node = values.next();
            Text outValue = new Text(nodeId.toString() + "\t\t" + node.toString());
            output.collect(key.getFirst(), outValue);
        }
    }
}
```

10. Program to print word count on scala shell and print “Hello world” on scala IDE

```
1 object HelloWorld {
2   def main(args: Array[String]) {
3     println("Hello world")
4   }
5 }
```

Hello world

HelloWorld.Scala (~)

File Edit View Search Tools Documents Help

⊕ 📁 💾 | ↶ ↷ | ✂ 📄 📋 | 🔍 🛠

HelloWorld.Scala ×

```
object HelloWorld
{
  def main(args: Array[String])
  {
    //This is a Hello World function in Scala
    println("Hello World!")
  }
}
```


11. Using RDD and FlatMap

Code:

```
>val text = sc.textFile("abc.txt")  
  
  >val counts = text.flatMap(line => line.split(" ")).map(word => (word,1)).reduceByKey(_+_)  
  counts.collect  
  
>val greaterThan4=counts.filter(x=>x._2>4);  
  
>greaterThan4.collect().foreach(println)
```

Input file:

Hello Hello World Hello Hello Xyz Xyz Xyz Hello World Hello Xyz World World Xyz Hello World

Output:

Hello 7

World 5

Xyz 5